

# **An introduction to tidy R programming with the OMOP common data model**

Edward Burn, Adam Black, Berta Raventós, Yuchen Guo, Mike Du, Kim López Güe

2023-02-13T00:00:00+00:00

# Table of contents


<b>Preface</b>	<b>3</b>
<b>1 Getting started</b>	<b>4</b>
1.1 A first data analysis in R with a database . . . . .	4
1.2 Getting set up . . . . .	4
1.3 Taking a peek at the data . . . . .	5
1.4 Inserting data into a database . . . . .	6
1.5 Translation from R to SQL . . . . .	7
1.6 Example analysis . . . . .	8
1.7 Further reading . . . . .	12
<b>2 Creating a reference to a database using the OMOP common data model</b>	<b>13</b>
2.1 Connecting to a database from R using DBI . . . . .	13
2.2 Creating a reference to the OMOP common data model . . . . .	13
2.3 Database snapshot . . . . .	16
2.4 Further reading . . . . .	16
<b>3 Exploring the CDM</b>	<b>17</b>
3.1 Counting people . . . . .	17
3.2 Counting records . . . . .	19
3.3 Working with dates . . . . .	20
3.4 Statistical summaries . . . . .	21
<b>4 Adding a cohort</b>	<b>24</b>
4.1 Defining a cohort definition . . . . .	24
4.2 Adding a cohort to the CDM . . . . .	24
4.3 Cohort summary . . . . .	24
<b>5 Describing a cohort</b>	<b>25</b>
5.1 CohortProfiles . . . . .	25
<b>References</b>	<b>26</b>

# Preface

The source code for the book can be found at this [Github repository](#). Please open an issue there if you have a question or suggestion. Pull requests with suggested changes and additions are also most welcome.

# 1 Getting started

## 1.1 A first data analysis in R with a database



*Artwork by @allison\_horst*

Before we start thinking about working with health care data spread across the OMOP common data model, let's first do a quick data analysis using a simpler dataset. For this we'll use data from [palmerpenguins package](#), which contains data on penguins collected from the [Palmer Station](#) in Antarctica.

## 1.2 Getting set up

Assuming that you have R and RStudio already set up, first we need to install a few packages not included in base R if we don't already have them.

```
install.packages("dplyr")
install.packages("ggplot2")
install.packages("DBI")
install.packages("duckdb")
install.packages("palmerpenguins")
```

Once installed, we can load them like so.

```
library(dplyr)
library(ggplot2)
library(DBI)
library(duckdb)
library(palmerpenguins)
```

## 1.3 Taking a peek at the data

We can get an overview of the data using the `glimpse()` command.

```
glimpse(penguins)
```

```
Rows: 344
Columns: 8
$ species      <fct> Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, Adel~
$ island       <fct> Torgersen, Torgersen, Torgersen, Torgersen, Torgerse~
$ bill_length_mm <dbl> 39.1, 39.5, 40.3, NA, 36.7, 39.3, 38.9, 39.2, 34.1, ~
$ bill_depth_mm <dbl> 18.7, 17.4, 18.0, NA, 19.3, 20.6, 17.8, 19.6, 18.1, ~
$ flipper_length_mm <int> 181, 186, 195, NA, 193, 190, 181, 195, 193, 190, 186~
$ body_mass_g   <int> 3750, 3800, 3250, NA, 3450, 3650, 3625, 4675, 3475, ~
$ sex          <fct> male, female, female, NA, female, male, female, male~
$ year         <int> 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007~
```

Or we could take a look at the first rows of the data using `head()`

```
head(penguins, 5)
```

```
# A tibble: 5 x 8
  species island  bill_length_mm bill_depth_mm flipper_l~1 body_~2 sex  year
  <fct>   <fct>      <dbl>         <dbl>      <int>    <int> <fct> <int>
1 Adelie Torgersen    39.1          18.7        181     3750 male   2007
2 Adelie Torgersen    39.5          17.4        186     3800 fema~  2007
3 Adelie Torgersen    40.3           18         195     3250 fema~  2007
4 Adelie Torgersen    NA           NA          NA        NA <NA>   2007
5 Adelie Torgersen    36.7          19.3        193     3450 fema~  2007
# ... with abbreviated variable names 1: flipper_length_mm, 2: body_mass_g
```

## 1.4 Inserting data into a database

Let's put our penguins data into a duckdb database. We create the duckdb database, add the penguins data, and then create a reference to the table containing the data.

```
db<-dbConnect(duckdb::duckdb(), dbdir=":memory:")
dbWriteTable(db, "penguins", penguins)
penguins_db<-tbl(db, "penguins")
```

Now the data is in a database we could use SQL to get the first rows that we saw before

```
dbGetQuery(db, "SELECT * FROM penguins LIMIT 5")
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g
1	Adelie	Torgersen	39.1	18.7	181	3750
2	Adelie	Torgersen	39.5	17.4	186	3800
3	Adelie	Torgersen	40.3	18.0	195	3250
4	Adelie	Torgersen	NA	NA	NA	NA
5	Adelie	Torgersen	36.7	19.3	193	3450

	sex	year
1	male	2007
2	female	2007
3	female	2007
4	<NA>	2007
5	female	2007

But we could also use the same R code as before

```
head(penguins_db, 5)
```

```
# Source:   SQL [5 x 8]
# Database: DuckDB 0.5.0 [eburn@Windows 10 x64:R 4.2.1/:memory:]
  species island  bill_length_mm bill_depth_mm flipper_l~1 body_~2 sex   year
  <fct>   <fct>         <dbl>         <dbl>         <int>   <int> <fct> <int>
1 Adelie  Torgersen      39.1           18.7           181     3750 male   2007
2 Adelie  Torgersen      39.5           17.4           186     3800 fema~  2007
3 Adelie  Torgersen      40.3            18            195     3250 fema~  2007
4 Adelie  Torgersen      NA              NA              NA        NA <NA>   2007
5 Adelie  Torgersen      36.7           19.3           193     3450 fema~  2007
# ... with abbreviated variable names 1: flipper_length_mm, 2: body_mass_g
```

## 1.5 Translation from R to SQL

The magic here is provided by dbplyr which takes the R code and converts it into SQL, which in this case looks like

```
head(penguins_db, 1) %>%  
  show_query()
```

```
<SQL>  
SELECT *  
FROM penguins  
LIMIT 1
```

More complicated SQL can also be written in what might be familiar dplyr code, for example

```
penguins_db %>%  
  group_by(species) %>%  
  summarise(min_bill_length_mm=min(bill_length_mm),  
            median_bill_length_mm=median(bill_length_mm),  
            max_bill_length_mm=max(bill_length_mm)) %>%  
  mutate(min_max_bill_length_mm=paste0(min_bill_length_mm,  
                                       " to ",  
                                       max_bill_length_mm)) %>%  
  select("species",  
        "median_bill_length_mm",  
        "min_max_bill_length_mm")
```

```
# Source:   SQL [3 x 3]  
# Database: DuckDB 0.5.0 [eburn@Windows 10 x64:R 4.2.1/:memory:]  
  species   median_bill_length_mm min_max_bill_length_mm  
  <fct>                <dbl> <chr>  
1 Adelie                38.8 32.1 to 46.0  
2 Gentoo                47.3 40.9 to 59.6  
3 Chinstrap            49.6 40.9 to 58.0
```

with the corresponding SQL looking like

```
penguins_db %>%  
  group_by(species) %>%  
  summarise(min_bill_length_mm=min(bill_length_mm),
```

```

        median_bill_length_mm=median(bill_length_mm),
        max_bill_length_mm=max(bill_length_mm)) %>%
mutate(min_max_bill_length_mm=paste0(min, " to ", max)) %>%
select("species",
       "median_bill_length_mm",
       "min_max_bill_length_mm") %>%
show_query()

```

<SQL>

```

SELECT
  species,
  median_bill_length_mm,
  CONCAT_WS(' ', .Primitive("min"), ' to ', .Primitive("max")) AS min_max_bill_length_mm
FROM (
  SELECT
    species,
    MIN(bill_length_mm) AS min_bill_length_mm,
    PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY bill_length_mm) AS median_bill_length_mm,
    MAX(bill_length_mm) AS max_bill_length_mm
  FROM penguins
  GROUP BY species
) q01

```

## 1.6 Example analysis

Let's start by getting a count by species

```

penguins_db %>%
  group_by(species) %>%
  count()

```

```

# Source:   SQL [3 x 2]
# Database: DuckDB 0.5.0 [eburn@Windows 10 x64:R 4.2.1/:memory:]
  species      n
  <fct>      <dbl>
1 Adelie     152
2 Gentoo     124
3 Chinstrap   68

```



Now suppose we are particularly interested in the body mass variable. We can first notice that there are a couple of missing records for this.

```
penguins_db %>%  
  mutate(missing_body_mass_g = if_else(  
    is.na(body_mass_g), 1, 0  
  )) %>%  
  group_by(species, missing_body_mass_g) %>%  
  tally()
```

```
# Source:   SQL [5 x 3]  
# Database: DuckDB 0.5.0 [eburn@Windows 10 x64:R 4.2.1/:memory:]  
# Groups:   species  
  species  missing_body_mass_g      n  
  <fct>                <dbl> <dbl>  
1 Adelie                0    151  
2 Adelie                1      1  
3 Gentoo                0    123  
4 Gentoo                1      1  
5 Chinstrap            0     68
```

We can get the mean for each of the species (dropping those two missing records).

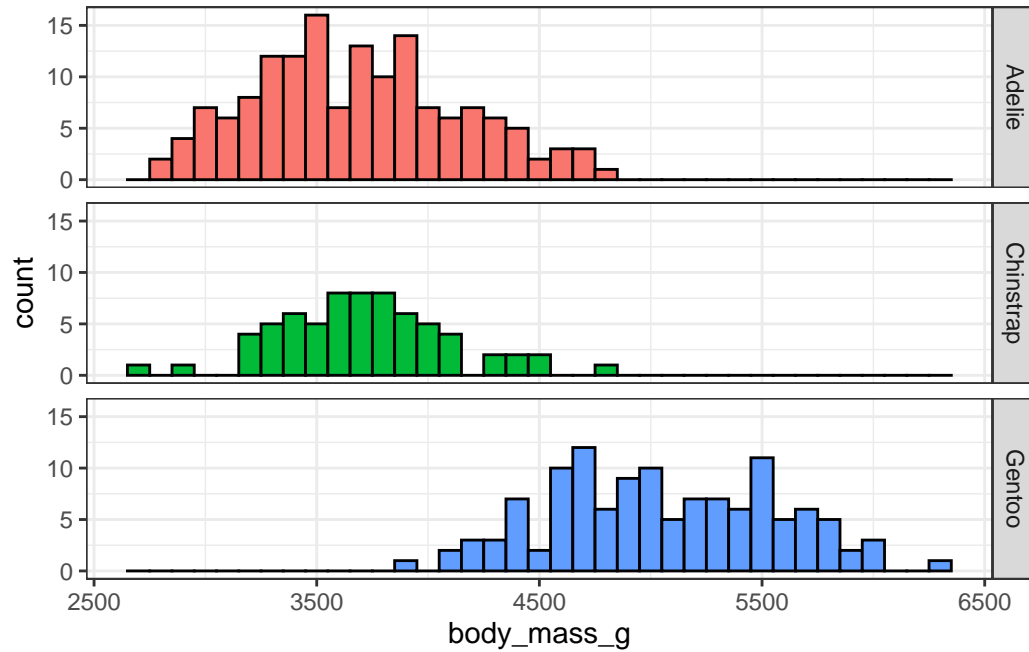
```
penguins_db %>%  
  group_by(species) %>%  
  summarise(mean_body_mass_g = round(mean(body_mass_g, na.rm=TRUE), 0))
```

```
# Source:   SQL [3 x 2]  
# Database: DuckDB 0.5.0 [eburn@Windows 10 x64:R 4.2.1/:memory:]  
  species  mean_body_mass_g  
  <fct>                <dbl>  
1 Adelie                3701  
2 Gentoo                5076  
3 Chinstrap            3733
```

We can then also do a histogram for each of the species. For this we need to bring the data into R so that we can work with `ggplot()`, and we use `collect()` to do this.


```
penguins_db %>%  
  collect() %>%
```

```
ggplot(aes(group=species, fill=species))+
  facet_grid(species~ .) +
  geom_histogram(aes(body_mass_g), colour="black", binwidth = 100)+
  theme_bw()+
  theme(legend.position = "none")
```




How about the relationship between body mass and bill depth?

```
penguins %>%
  collect() %>%
  ggplot(aes(x=bill_depth_mm,y=body_mass_g))+
  geom_point()+
  geom_smooth(method="lm",se=FALSE )+
  theme_bw()+
  theme(legend.position = "none")
```



But what about by species?

```
penguins %>%  
  collect() %>%  
  ggplot(aes(x=bill_depth_mm,y=body_mass_g))+  
  facet_grid(species~ .) +  
  geom_point()+  
  geom_smooth(method="lm",se=FALSE )+  
  theme_bw()+  
  theme(legend.position = "none")
```



As well as having an example of working with data in database from R, you also have an example of [Simpson's paradox](#)! And now we've reached the end of this example, we can close the database like so

## 1.7 Further reading

- [R for Data Science \(Chapter 13: Relational data\)](#)
- [Writing SQL with dbplyr](#)
- [Data Carpentry: SQL databases and R](#)

## 2 Creating a reference to a database using the OMOP common data model

### 2.1 Connecting to a database from R using DBI

Database connections from R can be made using the [DBI package](#). The back-end for DBI is facilitated by database specific driver packages. As an example, lets say we want to work with a local duckdb from R. In this case the we can use the duckdb R package as the driver.

```
library(DBI)
db<-dbConnect(duckdb::duckdb(), dbdir=":memory:")
```

If we instead wanted to connect to other database management systems, these connections would be supported by the associated back-end packages and could look something like the below example for Postgres:

```
# Postgres
db <- DBI::dbConnect(RPostgres::Postgres(),
                     dbname = Sys.getenv("CDM5_POSTGRESQL_DBNAME"),
                     host = Sys.getenv("CDM5_POSTGRESQL_HOST"),
                     user = Sys.getenv("CDM5_POSTGRESQL_USER"),
                     password = Sys.getenv("CDM5_POSTGRESQL_PASSWORD"))
```

### 2.2 Creating a reference to the OMOP common data model

As seen in the previous chapter, once a connection to the database has been created then we could create references to the various tables in the database and build queries using in a familiar dplyr style. However, as we already know what the structure of the OMOP CDM looks like, we can avoid the overhead of building *ad hoc* references by instead using the CDMConnector package to quickly create a reference to the OMOP CDM data as a whole.

If you don't already have it installed, the first step would be to install CDMConnector from CRAN.

```
install.packages("CDMConnector")
```

Once we have it installed, we can then load it as with other R packages.

```
library(CDMConnector)
```

For this example, we'll use the Eunomia example data contained in a duckdb database. First we need to download the data. And once downloaded, make sure to add the path to your Renviron.

```
# change pathToData to the location you want to save the data
CDMConnector::downloadEunomiaData(
  pathToData = here::here(),
  overwrite = TRUE
)
# once downloaded, save your pathToData to your Renviron (and then restart R)
# EUNOMIA_DATA_FOLDER="....."
```

```
db <- DBI::dbConnect(duckdb::duckdb(),
  dbdir = CDMConnector::eunomia_dir())
cdm <- CDMConnector::cdm_from_con(con = db,
  cdm_schema = "main")

cdm
```

```
# OMOP CDM reference (tbl_duckdb_connection)
```

Tables: person, observation\_period, visit\_occurrence, visit\_detail, condition\_occurrence, drug\_exposure

Once we have created the our reference to the overall OMOP CDM, we can reference specific tables using the “\$” operator or [[ “”]].

```
cdm$observation_period
```

```
# Source:   table<main.observation_period> [?? x 5]
# Database: DuckDB 0.5.0 [eburn@Windows 10 x64:R 4.2.1/C:\Users\eburn\AppData\Local\Temp\Rtmp]
  observation_period_id person_id observation_period_start~1 observat~2 perio~3
      <dbl>          <dbl> <date>                  <date>          <dbl>
1           6           6 1963-12-31          2007-02-06  4.48e7
2          13          13 2009-04-26          2019-04-14  4.48e7
3          27          27 2002-01-30          2018-11-21  4.48e7
```

```

4          16          16 1971-10-14          2017-11-02 4.48e7
5          55          55 2009-05-30          2019-03-23 4.48e7
6          60          60 1990-11-21          2019-01-23 4.48e7
7          42          42 1909-11-03          2019-03-13 4.48e7
8          33          33 1986-05-12          2018-09-10 4.48e7
9          18          18 1965-11-17          2018-11-07 4.48e7
10         25          25 2007-03-18          2019-04-07 4.48e7
# ... with more rows, and abbreviated variable names
#   1: observation_period_start_date, 2: observation_period_end_date,
#   3: period_type_concept_id

```

```
cdm[["observation_period"]]
```

```

# Source:   table<main.observation_period> [?? x 5]
# Database: DuckDB 0.5.0 [eburn@Windows 10 x64:R 4.2.1/C:\Users\eburn\AppData\Local\Temp\Rtmp
  observation_period_id person_id observation_period_start~1 observat~2 perio~3
                <dbl>      <dbl> <date>                <date>      <dbl>
1              6         6 1963-12-31          2007-02-06 4.48e7
2             13        13 2009-04-26          2019-04-14 4.48e7
3             27        27 2002-01-30          2018-11-21 4.48e7
4             16        16 1971-10-14          2017-11-02 4.48e7
5             55        55 2009-05-30          2019-03-23 4.48e7
6             60        60 1990-11-21          2019-01-23 4.48e7
7             42        42 1909-11-03          2019-03-13 4.48e7
8             33        33 1986-05-12          2018-09-10 4.48e7
9             18        18 1965-11-17          2018-11-07 4.48e7
10            25        25 2007-03-18          2019-04-07 4.48e7
# ... with more rows, and abbreviated variable names
#   1: observation_period_start_date, 2: observation_period_end_date,
#   3: period_type_concept_id

```

When we created our reference we could have also specified a subset of cdm tables that we want to read:

```

cdm <- CDMConnector::cdm_from_con(db,
                                   cdm_tables = c("person","observation_period"))
cdm

```

```
# OMOP CDM reference (tbl_duckdb_connection)
```

```
Tables: person, observation_period
```

Moreover, we can also specify a write schema and the tables that we are interested in it when creating our reference. For example, if we wanted to create a reference to the person and observation period tables in the common data model along with cohort tables in a schema we have write access to, we could do this like so:

```
cdm <- CDMConnector::cdm_from_con(db,
  cdm_schema = "main",
  cdm_tables = c("person", "observation_period"),
  write_schema = "results",
  cohort_tables = c("exposure_cohort", "outcome_cohort"))
```

## 2.3 Database snapshot

We can also use `CDMConnector` to provide a summary of the metadata for the OMOP CDM data we have connected to

```
cdm_from_con(con = db,
  cdm_schema = "main") %>%
  snapshot() %>%
  glimpse()
```

List of 7

```
$ cdm_source_name      : chr "Synthea synthetic health database"
$ cdm_version          : chr "v5.3.1"
$ cdm_holder           : chr "OHDSI Community"
$ cdm_release_date     : Date[1:1], format: "2019-05-25"
$ vocabulary_version   : chr "v5.0 18-JAN-19"
$ person_cnt           : num 2694
$ observation_period_cnt: num 5343
- attr(*, "class")= chr "cdm_snapshot"
```

## 2.4 Further reading

- [CDMConnector package](#)



## 3 Exploring the CDM

Let's first connect again to our Eunomia data and create the reference to the common data model.

```
library(dbplyr)
library(dplyr)
library(CDMConnector)
library(ggplot2)
```

### 3.1 Counting people

The OMOP CDM is person-centric, with the person table containing records to uniquely identify each person in the database. As each row refers to a unique person, we can quickly get a count of the number of individuals in the database like so

```
cdm$person %>%
  count() %>%
  pull()
```

```
[1] 2694
```

The person table also contains some demographic information, including a `gender_concept_id` for each person. We can get a count grouped by this variable, but as this uses a concept we'll also need to join to the concept table to get the corresponding concept name for each concept id.

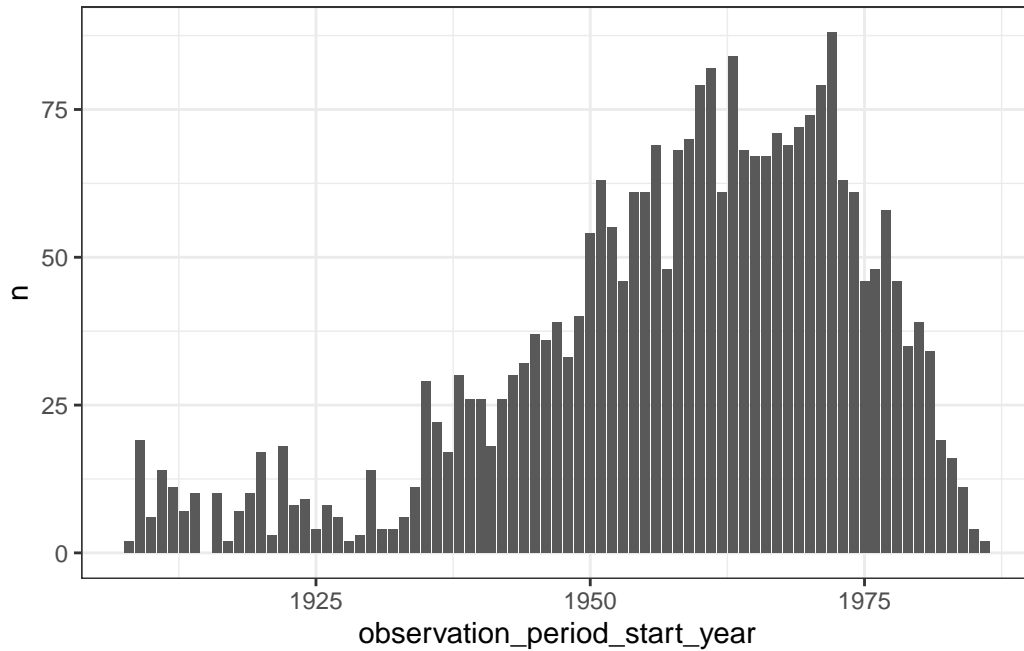
```
cdm$person %>%
  group_by(gender_concept_id) %>%
  count() %>%
  left_join(cdm$concept,
            by=c("gender_concept_id" = "concept_id")) %>%
    select("gender_concept_id", "concept_name", "n") %>%
  collect()
```

```
# A tibble: 2 x 3
  gender_concept_id concept_name      n
      <dbl> <chr>      <dbl>
1         8532 FEMALE      1373
2         8507 MALE       1321
```

The observation period table contains records indicating spans of time over which clinical events can be reliably observed for the people in the person table. Someone can potentially have multiple observation periods. So say we wanted a count of people grouped by the year during which their first observation period started. We could do this as below (note the use of `compute()` to store the results of the first query in a temporary table in the database)

```
first_observation_period <- cdm$observation_period %>%
  group_by(person_id) %>%
  mutate(seq = dplyr::row_number()) %>%
  filter(seq==1) %>%
  compute()


cdm$person %>%
  left_join(first_observation_period,
            by = "person_id") %>%
  mutate(observation_period_start_year=year(observation_period_start_date)) %>%
  group_by(observation_period_start_year) %>%
  count() %>%
  collect() %>%
  ggplot() +
  geom_col(aes(observation_period_start_year, n)) +
  theme_bw()
```



## 3.2 Counting records

Number of drug exposure records per person

```
cdm$person %>%
  left_join(cdm$measurement %>%
    group_by(person_id) %>%
    count(name = "condition_occurrence_records",
          by="person_id") %>%
    mutate(condition_occurrence_records = if_else(
      is.na(condition_occurrence_records), 0,
      condition_occurrence_records)) %>%
    group_by(condition_occurrence_records) %>%
    count() %>%
    collect() %>%
  ggplot() +
  geom_col(aes(condition_occurrence_records, n)) +
  theme_bw()
```




### 3.3 Working with dates

Dates are supported somewhat inconsistently by dbplyr, but CDMConnector provides some functions that provide more general support. We can use the `datediff` function from CDMConnector for example to calculate the difference between two dates. We can use this, for example, to get the number of years people's observation period last for.

```
cdm$observation_period %>%
  dplyr::mutate(observation_years =
    !!CDMConnector::datediff("observation_period_start_date",
                             "observation_period_end_date",
                             interval = "year")) %>%

  collect() %>%
  ggplot() +
  geom_histogram(aes(observation_years),
                 binwidth=2, colour="grey") +
  theme_bw()
```



### 3.4 Statistical summaries

We can also use summarise for various other calculations

```
cdm$person %>%
  summarise(min_year_of_birth = min(year_of_birth, na.rm=TRUE),
            q05_year_of_birth = quantile(year_of_birth, 0.05, na.rm=TRUE),
            mean_year_of_birth = round(mean(year_of_birth, na.rm=TRUE),0),
            median_year_of_birth = median(year_of_birth, na.rm=TRUE),
            q95_year_of_birth = quantile(year_of_birth, 0.95, na.rm=TRUE),
            max_year_of_birth = max(year_of_birth, na.rm=TRUE)) %>%
  glimpse()
```

Rows: ??

Columns: 6

Database: DuckDB 0.5.0 [eburn@Windows 10 x64:R 4.2.1/C:\Users\eburn\AppData\Local\Temp\RtmpA

\$ min\_year\_of\_birth <dbl> 1908

\$ q05\_year\_of\_birth <dbl> 1922

\$ mean\_year\_of\_birth <dbl> 1958


\$ median\_year\_of\_birth <dbl> 1961

\$ q95\_year\_of\_birth <dbl> 1979

```
$ max_year_of_birth    <dbl> 1986
```

As we've seen before, we can also quickly get results for various groupings or restrictions

```
cdm$person %>%
  group_by(gender_concept_id) %>%
  summarise(min_year_of_birth = min(year_of_birth, na.rm=TRUE),
            q25_year_of_birth = quantile(year_of_birth, 0.25, na.rm=TRUE),
            median_year_of_birth = median(year_of_birth, na.rm=TRUE),
            q75_year_of_birth = quantile(year_of_birth, 0.75, na.rm=TRUE),
            max_year_of_birth = max(year_of_birth, na.rm=TRUE)) %>%
  left_join(cdm$concept,
            by=c("gender_concept_id" = "concept_id")) %>%
  collect() %>%
  ggplot(aes(x = concept_name, group = concept_name,
            fill = concept_name)) +
  geom_boxplot(aes(
    lower = q25_year_of_birth,
    upper = q75_year_of_birth,
    middle = median_year_of_birth,
    ymin = min_year_of_birth,
    ymax = max_year_of_birth),
    stat = "identity", width = 0.5) +
  theme_bw()+
  theme(legend.position = "none") +
  xlab("")
```



## **4 Adding a cohort**

### **4.1 Defining a cohort definition**

capr - defining json ....

### **4.2 Adding a cohort to the CDM**

via cdm connector

### **4.3 Cohort summary**

Attributes: Cohort count, settings, attrition .....



## **5 Describing a cohort**

### **5.1 CohortProfiles**

### **5.2**

## References